

Aula 3: Treino e inferência Transformer Decoder-Only

Prof. Marcelo Keese Albertini

Universidade Federal de Uberlândia



Roteiro

1. Revisão Transformer
2. Transformer Decoder-Only
3. Embedding
4. Codificação de posições
5. Inferência
6. Prática
7. Referências

Transformer

“Attention Is All You Need” (2017)

<https://arxiv.org/abs/1706.03762>




Transformer Decoder-Only

Texto → `Tokenizer.encode` → Tokens →

Modelo:

1. Embedding + Positional Encoding (PE) →
2. [Multi-head **Causal** Self-Attention → FFN] × N →
3. Linear →
4. Softmax →

Inferência do próximo token → `Tokenizer.decode` → Texto



Embedding


```
>>> # an Embedding module containing 10 tensors of size 3
>>> embedding = nn.Embedding(10, 3)
>>> # a batch of 2 samples of 4 indices each
>>> input = torch.LongTensor([[1, 2, 4, 5], [4, 3, 2, 9]])
>>> embedding(input)
tensor([[[[-0.0251, -1.6902,  0.7172],
          [-0.6431,  0.0748,  0.6969],
          [ 1.4970,  1.3448, -0.9685],
          [-0.3677, -2.7265, -0.1685]],

         [[ 1.4970,  1.3448, -0.9685],
          [ 0.4362, -0.4004,  0.9400],
          [-0.6431,  0.0748,  0.6969],
          [ 0.9124, -2.3616,  1.1151]]])
```

<https://docs.pytorch.org/docs/stable/generated/torch.nn.Embedding.html>



<https://zackproser.com/demos/embeddings>




Codificação de posições: Positional Encoding (PE)

$$\text{Input} = \text{Embedding}(token) + \text{PE}(pos)$$

- **Attention** não tem como representar a ordem dos tokens por si só !
- **Positional Encoding** demarca a posição dos tokens nas frases ao longo das dimensões do **Embedding**.

https://colab.research.google.com/drive/1gvdxzcMYQwooPaSH4_ifh5Wc8Sn6CXJw?usp=sharing



Positional Encoding (PE)


$$\text{Input} = \text{Embedding}(token) + \text{PE}(pos)$$

$$\text{PE}_{(pos,2i)} = \sin\left(\frac{pos}{10000^{2i/d_{\text{model}}}}\right)$$

$$\text{PE}_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{2i/d_{\text{model}}}}\right)$$


- **pos**: posição do token na sequência de tokens
- **i**: dimensão do embedding
- **d_model**: número total de dimensões no embedding

```
def positional_encoding(pos, i, d_model):
    angle = pos / (10000 ** (2 * (i//2) / d_model))
    return np.sin(angle) if i % 2 == 0 else np.cos(angle)
```



Positional Encoding (PE)

- PE aumenta similaridade entre time-steps próximos
 - aumento de **similaridade** entre o time 0 e o 50 é **baixo**
 - aumento de **similaridade** entre o time-step 49 e 50 é **alto**



https://kazemnejad.com/blog/transformer_architecture_positional_encoding/

Figure 3 - Dot product of position embeddings for all time-steps

Inferência: decoding - generate

```
def generate(prompt, next_token_function, max_new_tokens=20):
    model.eval()
    x = encode(prompt).unsqueeze(0).to(device) # tokenizer

    for _ in range(max_new_tokens):
        logits = model(x)[:, -1, :] # pega apenas o último passo
        next_token = next_token_function(logits.squeeze(0))

        x = torch.cat([x, next_token.unsqueeze(0)], dim=1)

        if next_token.item() == tokenizer.token_to_id("[EOS]"):
            break

    return decode(x[0])
```

```
def max_prob_sampling(logits):
    next_token = logits.argmax(dim=-1)
    return next_token.unsqueeze(0)
```

Inferência: decoding - estratégias de amostra de token

```
def sampling(logits, top_k=100, top_p=0.9, temperature=1.0):
```


1. **Logits**: probabilidades dos tokens produzidas pelo modelo

2. **Aplicar Temperatura**: modificador de probabilidades.

Se > 1 , aumenta chance de tokens menos prováveis e mais **criativo**. Se < 1 , deixa mais **rígido**.

3. **Filtro dos Top-k**: Somente os tokens com as **k** maiores probabilidades são considerados

4. **Amostragem Top-p**: exclui os tokens cuja probabilidade acumulada não inclui p%.



Controles: temperatura, top-k, top-P

Visualização de controle de Temperatura, Top-K e Top-P.

Parameters

Prompt ▼

Temperatura 1.5

Top-K 6

Top-P 1.00



<https://colab.research.google.com/drive/16g7RjiELBvtveKQ7hHxaKgZVAHwDbAH3?usp=sharing>

<https://andreban.github.io/temperature-topk-visualizer>

Referências

- Como inferência de LLMs funciona
 - <https://bentoml.com/llm/llm-inference-basics/how-does-llm-inference-work>
- Deep dive into Text Generation Inference with LLMs
 - <https://huggingface.co/learn/llm-course/en/chapter1/8>
- CSCI1470 - Deep learning
 - <https://deep-learning-s25.vercel.app/>
- “The Annotated Transformer” (Harvard NLP)
 - <https://nlp.seas.harvard.edu/2018/04/03/attention.html>
- “Transformers from Scratch” — Peter Bloem
 - <https://peterbloem.nl/blog/transformers>
- Positional encoding
 - https://kazemnejad.com/blog/transformer_architecture_positional_encoding/